

Spatial graphics

Rolf Bänziger, r.banziger@westminster.ac.uk

24 November 2022

Spatial (or geospatial) data is data related to specific locations. We often want to display this kind of data overlaid on a map. This tutorial demonstrates producing geospatial graphics with *ggplot2*.

Libraries

We need to install a number of libraries first. *sf* provides functions for working with geospatial data. See [Simple Feature for R](#) for more details. *rnatuarearth* and *rnatuarearthdata* provide a world map with some demographics data, and *eurostat* provides a number of statistics from the European Union. We also need *tidyverse*, which we use to manipulate data.

```
install.packages(c("sf", "rnaturalearth", "rnaturalearthdata", "rgeos",
                   "eurostat", "tidyverse", "cartogram", "s2"))
```

```
library(ggplot2)
library(sf)
```

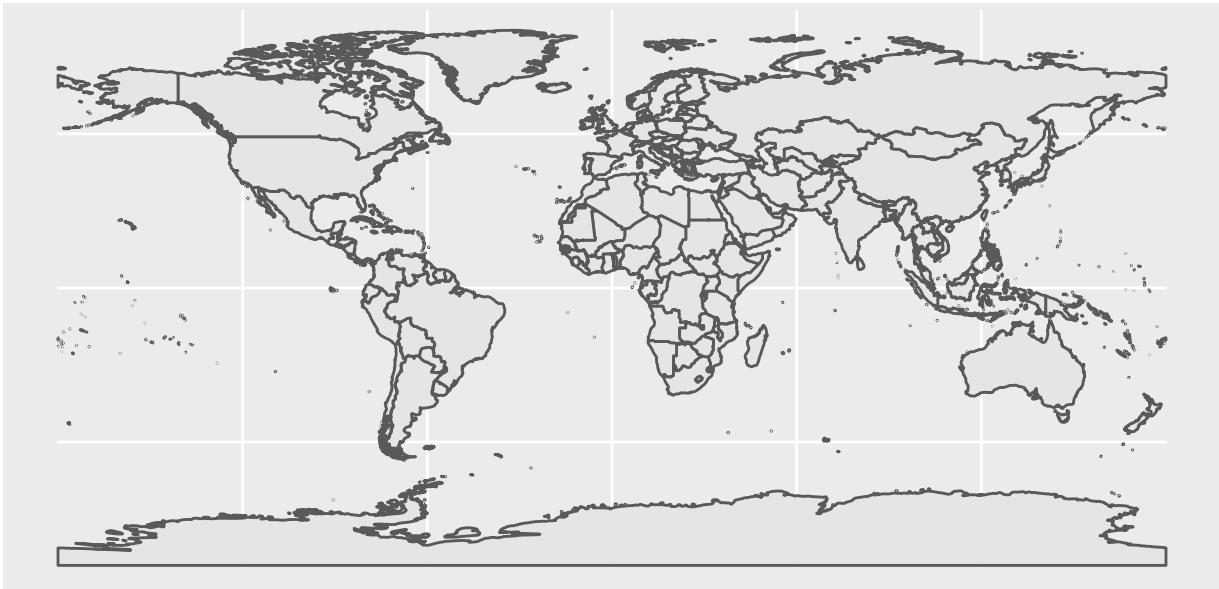
```
## Warning: package 'sf' was built under R version 4.2.2
```

```
library(rnaturalearth)
```

```
## Warning: package 'rnaturalearth' was built under R version 4.2.2
```

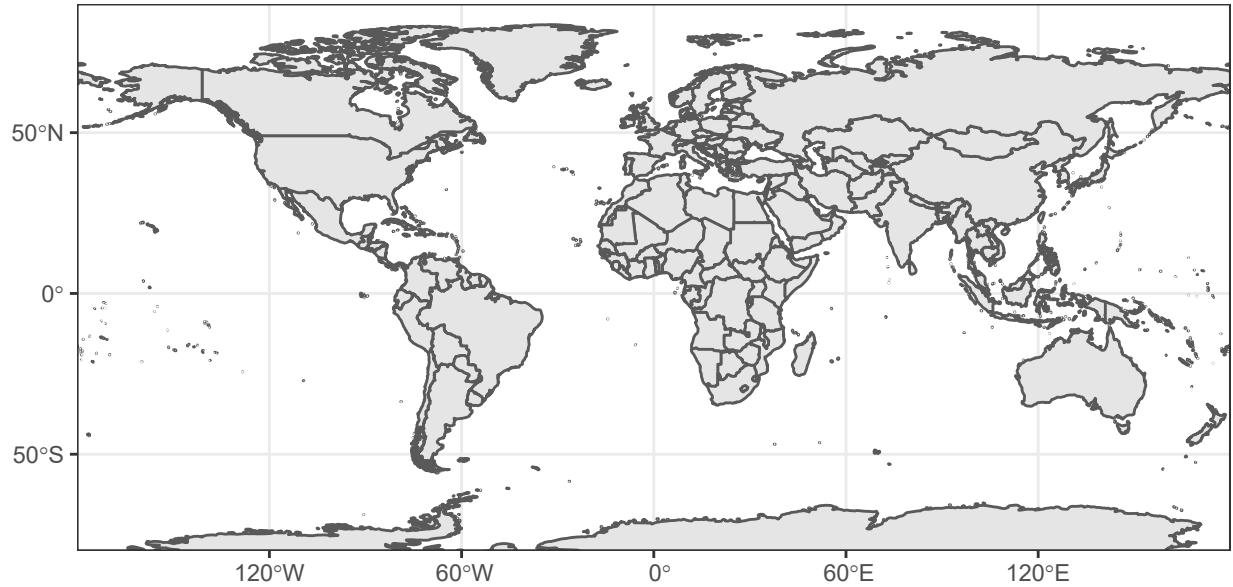
Let's get the world map and save it in the variable `world`. Now we can plot it using `geom_sf()`. Note that we don't need to pass any aesthetics mappings just yet.

```
world <- ne_countries(returnclass = "sf", scale="medium")
ggplot(world) + geom_sf()
```



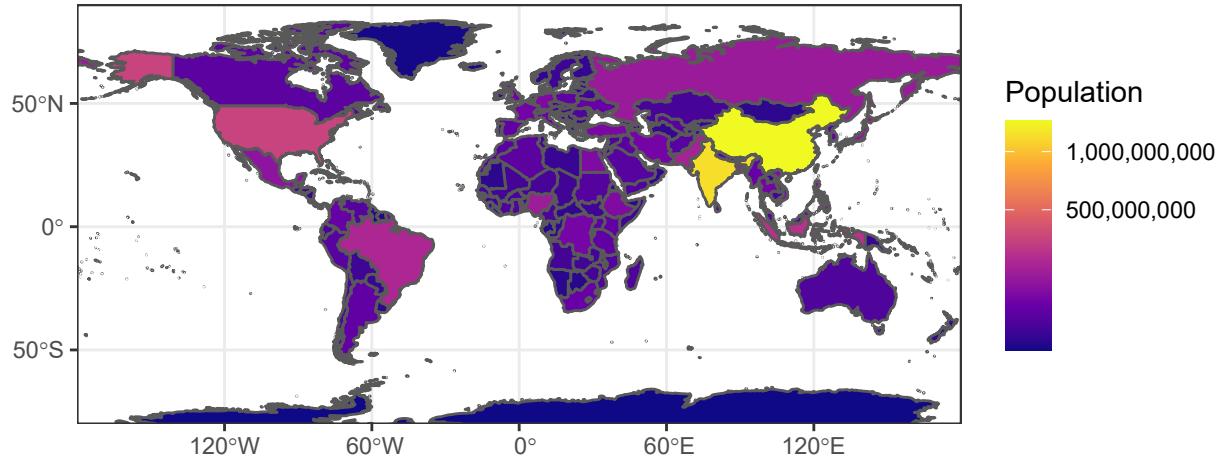
Let's make it a bit nicer. `theme_bw()` works well with maps and `coord_sf(expand = F)` ensures the latitude and longitude are added. We also specify northern and southern limits.

```
ggplot(world) + geom_sf() +
  theme_bw() + coord_sf(ylim = c(-80, 90), expand = F)
```



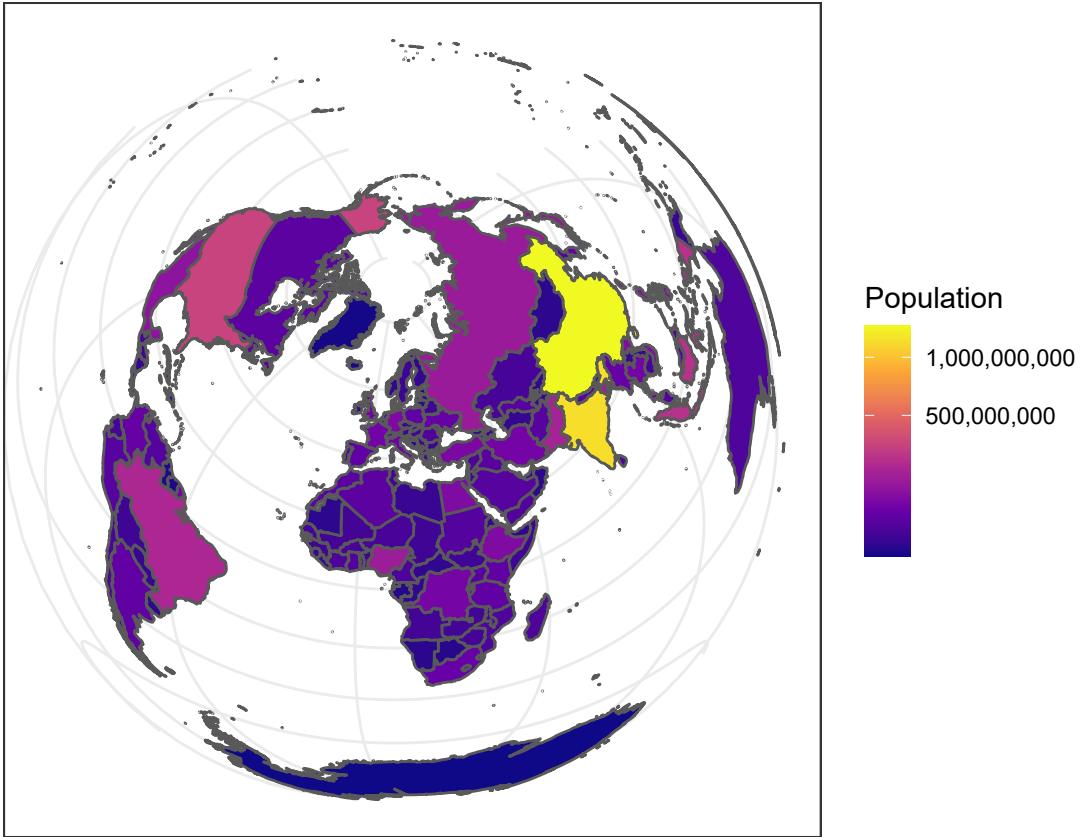
`geom_sf()` supports mapping data to the *fill* aesthetics. Let's use this to display the estimated population. Note that we transform the scale into an exponential scale by adding the `trans = "sqrt"` option. Otherwise the scale would be completely dominated by China and India.

```
ggplot(world) + geom_sf() + aes(fill = pop_est) +  
  scale_fill_viridis_c(option = "plasma", trans = "sqrt", label = scales::comma) +  
  labs(fill = "Population") +  
  theme_bw() + coord_sf(ylim = c(-80, 90), expand = F)
```



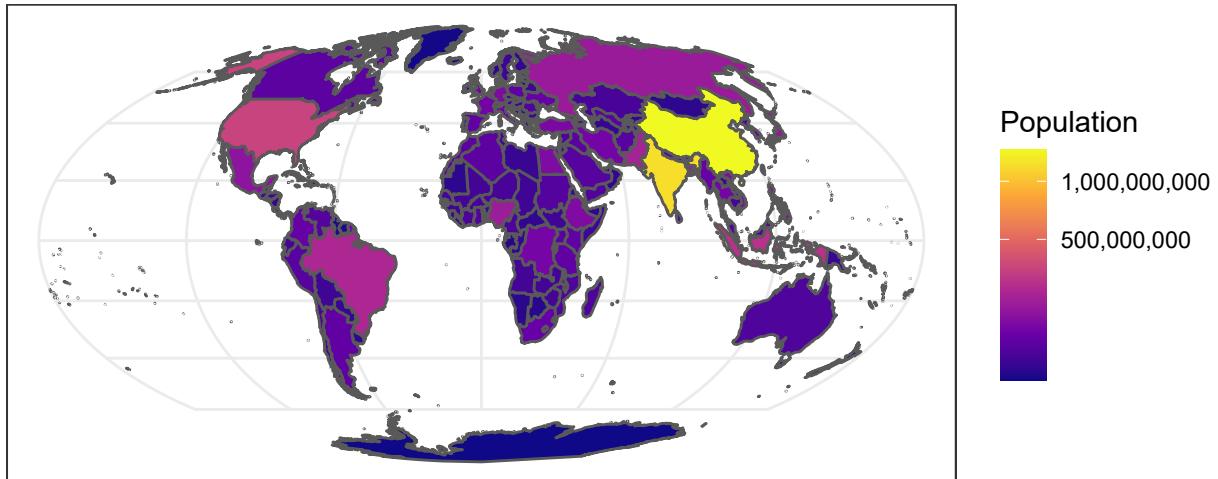
The function `coord_sf` allows setting the coordinate reference system, which includes both projection and extent of the map. By default, the map will use the coordinate system that is defined within the map data or, if none is defined, WGS84 (latitude/longitude, the reference system used in GPS). Use the argument `crs` to override this setting. Here, we set it the EPSG (European Petroleum Survey Group) code 3035, which is the European-centric ETRS89 (European Terrestrial Reference System 1989) Lambert Azimuthal Equal-Area projection. There are many other EPSG codes, see [EPSG.org](#) for a list of all available codes.

```
ggplot(world) + geom_sf() + aes(fill = pop_est) +
  scale_fill_viridis_c(option = "plasma", trans = "sqrt", label = scales::comma) +
  labs(fill = "Population") +
  theme_bw() +
  coord_sf(crs = st_crs(3035))
```



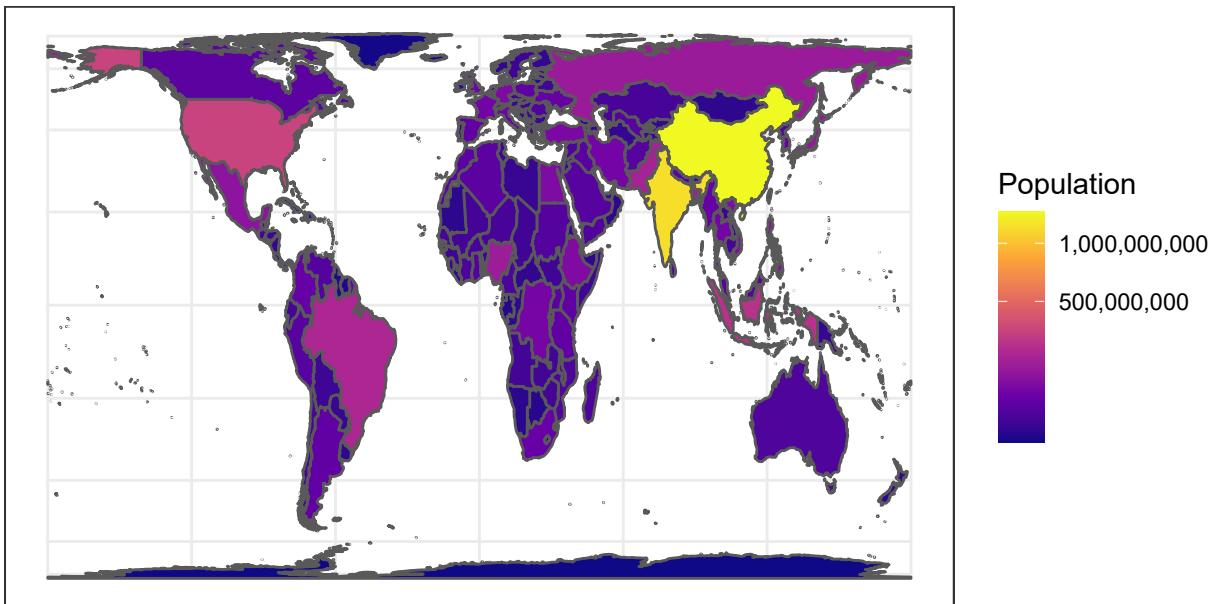
Alternativley, you can directly specify the projection using a *proj4string*. See [Proj4](#) for a list of available projections. The example below shows the *Mollweide* projection:

```
ggplot(world) + geom_sf() + aes(fill = pop_est) +  
  scale_fill_viridis_c(option = "plasma", trans = "sqrt", label = scales::comma) +  
  labs(fill = "Population") +  
  theme_bw() +  
  coord_sf(crs = "+proj=moll")
```



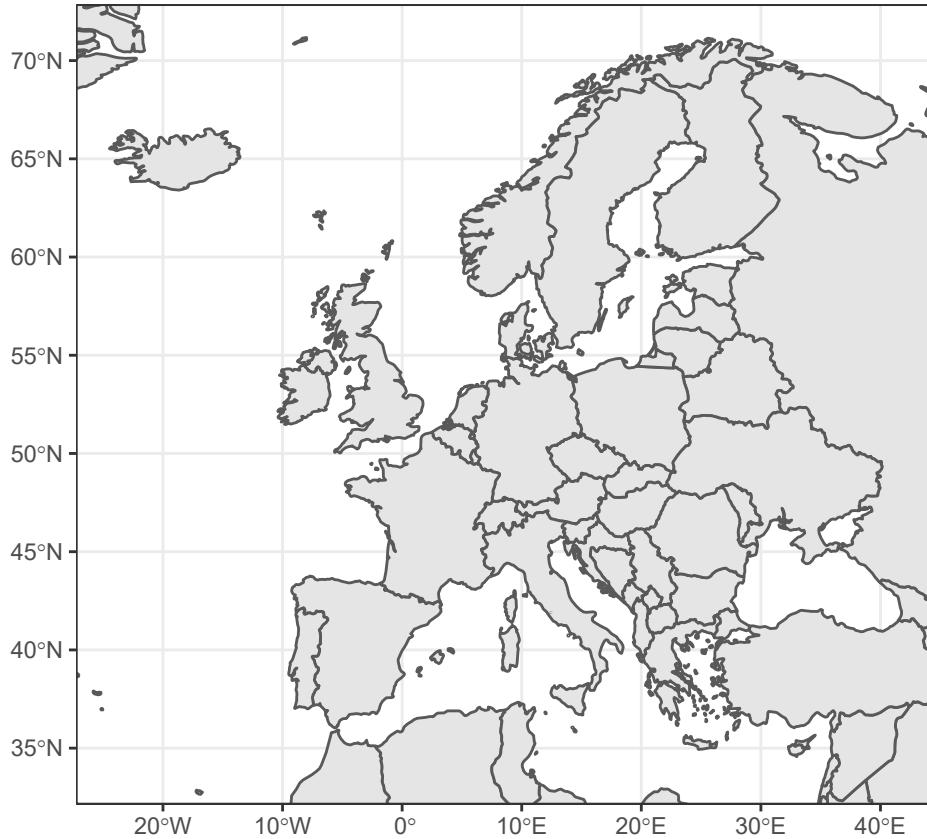
And this map uses the Gall-Peters projection:

```
ggplot(world) + geom_sf() + aes(fill = pop_est) +  
  scale_fill_viridis_c(option = "plasma", trans = "sqrt", label = scales::comma) +  
  labs(fill = "Population") +  
  theme_bw() +  
  coord_sf(crs = "+proj=cea +lon_0=0 +lat_ts=45")
```



We can also use the `coord_sf` function to zoom into a specific region, by specifying latitude and longitude ranges. Note that the latitude relates to the y -axis and the longitude to the x-axis. Negative values refer to southern or western coordinates, whereas positive values refer to northern or eastern coordinates.

```
ggplot(world) + geom_sf() + theme_bw() +  
  coord_sf(ylim = c(34, 71), xlim = c(-24, 41))
```



Europe

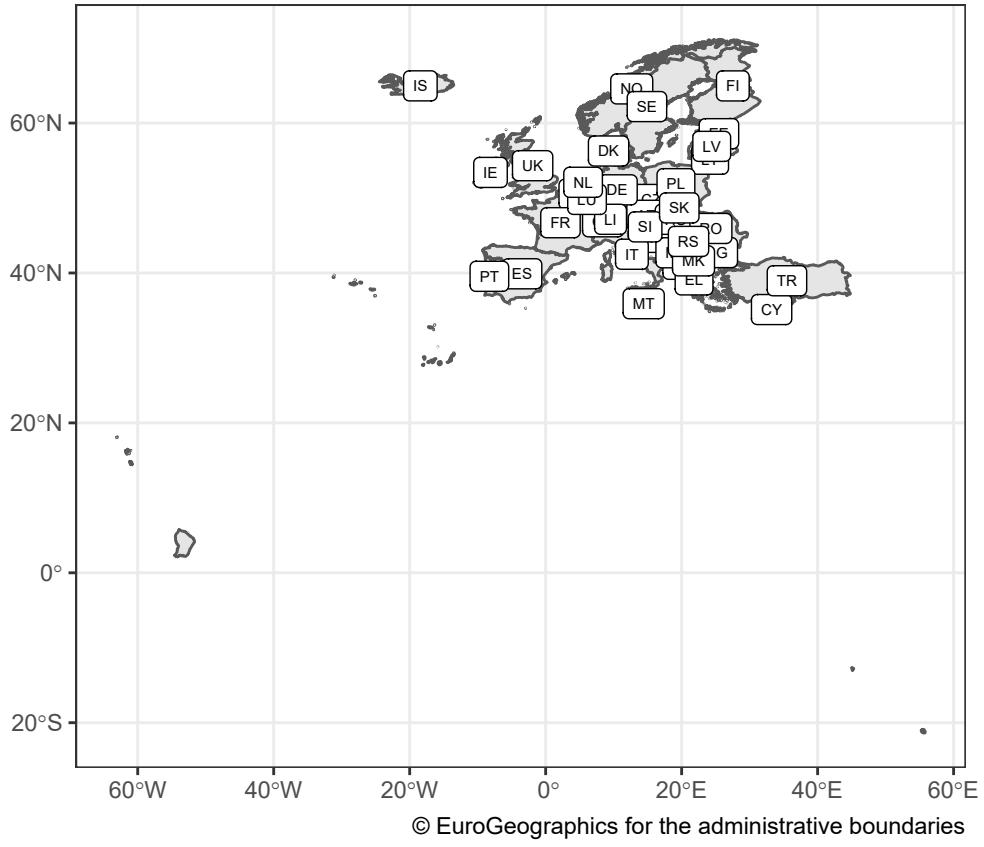
Let's draw some maps with European data. The package *eurostat* provides functions to get maps and data from [Eurostat](#).

```
library(eurostat)

## Warning: package 'eurostat' was built under R version 4.2.2
```

We can download the map with national boundaries using `get_eurostat_geospatial(nuts = 0)`. See [NUTS](#) for details about the regional levels. As usual, we use `ggplot()` and `geom_sf()` to draw the map. We also draw labels for each country with `geom_sf_text()`. Because the map data is provided by Eurostat, we need to add a copyright note.

```
eurostat.map <- get_eurostat_geospatial(nuts = 0)
ggplot(eurostat.map) + geom_sf() +
  geom_sf_label(aes(label = id), size = 2) +
  theme_bw() +
  labs(caption = "© EuroGeographics for the administrative boundaries",
       x = NULL, y = NULL)
```



Oops, looks like the map includes some colonies, let's crop it to only show continental europe...

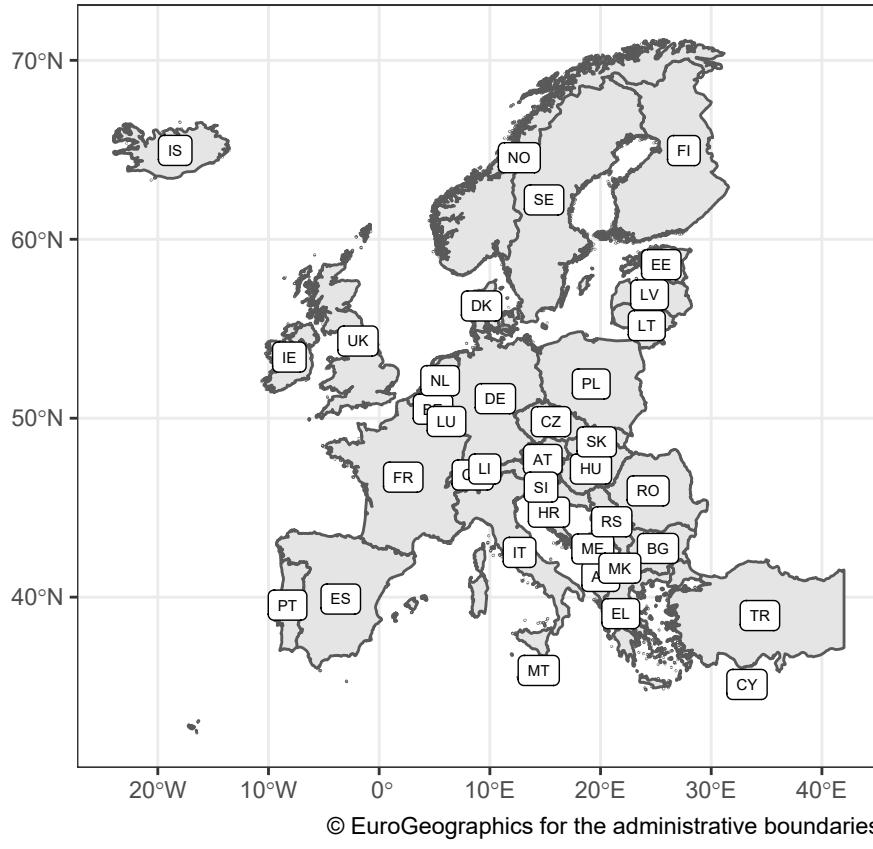
```
eurostat.map <- st_crop(eurostat.map, xmin = -24, xmax = 42,
                           ymin = 35, ymax = 70)

## Warning: attribute variables are assumed to be spatially constant throughout all
## geometries

ggplot(eurostat.map) +geom_sf() +
  geom_sf_label(aes(label = id), size = 2) +
  theme_bw()

  labs(caption = "© EuroGeographics for the administrative boundaries",
       x = NULL, y = NULL)

## Warning in st_point_on_surface.sfc(sf::st_zm(x)): st_point_on_surface may not
## give correct results for longitude/latitude data
```



Let us load the demographic at national level (`demo_gind`). Consult the [Eurostat Database](#) for a list of all available data sets.

```
demographics <- get_eurostat("demo_gind")
head(demographics)
```

```
## # A tibble: 6 x 4
##   indic_de geo    time      values
##   <chr>    <chr> <date>     <dbl>
## 1 JAN      AD    2022-01-01    79535
## 2 JAN      AL    2022-01-01  2793592
## 3 JAN      AT    2022-01-01  8978929
## 4 JAN      AZ    2022-01-01 10156366
## 5 JAN      BE    2022-01-01 11631136
## 6 JAN      BG    2022-01-01  6838937
```

We see that the data set above contains one row per country, statistical index and year. We'd rather have only the data from 2020, and only one row per country containing all indices. Let's use `tidyverse` functions to filter and pivot the data.

```
library(tidyverse)

## Warning: package 'tidyverse' was built under R version 4.2.2
```

```

demo2 <-
  demographics %>%
  filter(time == as.Date("2020-01-01")) %>%
  pivot_wider(names_from = indic_de, values_from = values)
head(demo2)

```

```

## # A tibble: 6 x 30
##   geo    time          AVG CNMIGRAT CNMIGRATRT DEATH     FAVG FDEATH    FJAN
##   <chr> <date>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 AL    2020-01-01  2837848   -16684     -5.9    27605  1422551  12101  1425342
## 2 AM    2020-01-01  2961473    3374      1.1     36170  1563917  16139  1562689
## 3 AT    2020-01-01  8916864   39596      4.4     91599  4529002  46227  4522292
## 4 AZ    2020-01-01  10093121   1101      0.1     75647  5052194  33566  5039100
## 5 BE    2020-01-01  11538604   44873      3.9    126896  5847754  64688  5841215
## 6 BG    2020-01-01  6934015   30715      4.4    124735  3574335  58785  3581836
## # ... with 21 more variables: FLBIRTH <dbl>, GBIRTHRT <dbl>, GDEATHRT <dbl>,
## #   GROW <dbl>, GROWRT <dbl>, JAN <dbl>, LBIRTH <dbl>, MAVG <dbl>,
## #   MDEATH <dbl>, MIGT <dbl>, MIGTRT <dbl>, MJAN <dbl>, MLBIRTH <dbl>,
## #   NATGROW <dbl>, NATGROWRT <dbl>, NATTRT <dbl>, POPSHARE <dbl>,
## #   POPSHARE_EU27_2020 <dbl>, POPT <dbl>, POPTRT <dbl>

```

That's looks much better. To use the data in our geo-plot, we need to join the map and the demo2 data frames. We use the `left_join` function from tidyverse, which will retain all values from `eurostat.map` and only use matching records from `demo2`.

```

demomap <- left_join(eurostat.map, demo2, by = "geo")
head(as.data.frame(demomap))

```

```

##   id CNTR_CODE          NUTS_NAME LEVL_CODE FID NUTS_ID geo      time
## 1 BG   BG              0         BG   BG   BG 2020-01-01
## 2 CH   CH  SCHWEIZ/SUISSE/SVIZZERA 0         CH   CH 2020-01-01
## 3 CY   CY              KΤΠΡΟΣ 0         CY   CY 2020-01-01
## 4 AL   AL              SHQIPËRIA 0         AL   AL 2020-01-01
## 5 CZ   CZ              ČESKÁ REPUBLIKA 0         CZ   CZ 2020-01-01
## 6 BE   BE              BELGIQUE-BELGIË 0         BE   BE 2020-01-01
##   AVG CNMIGRAT CNMIGRATRT DEATH     FAVG FDEATH    FJAN FLBIRTH GBIRTHRT
## 1 6934015 30715        4.4 124735 3574335 58785 3581836 28714     8.5
## 2 8638167 54548        6.3 76195 4352436 38570 4337170 41615     9.9
## 3 892006  4494         5.0 6422  455946 3038  453534  4740    11.1
## 4 2837848 -16684       -5.9 27605 1422551 12101 1425342 13589     9.9
## 5 10697858 26927        2.5 129289 5424309 62690 5421943 53924    10.3
## 6 11538604 44873        3.9 126896 5847754 64688 5841215 55843     9.9
##   GDEATHRT GROW GROWRT      JAN LBIRTH     MAVG MDEATH    MIGT MIGTRT    MJAN
## 1 18.0 -34934 -5.0 6951482 59086 3359681 65950 44013  6.3 3369646
## 2 8.8  64267  7.4 8606033 85914 4285731 37625 267528 31.0 4268863
## 3 7.2  8002   9.0 888005 9930  436061 3384  47230 52.9 434471
## 4 9.7 -16214 -5.7 2845955 28075 1415298 15504 31024 10.9 1420613
## 5 12.1  7838   0.7 10693939 110200 5273550 66599 98635 9.2 5271996
## 6 11.0  32327  2.8 11522440 114350 5690850 62208 182727 15.8 5681225
##   MLBIRTH NATGROW NATGROWRT    NATT NATTRT POPSHARE POPSHARE_EU27_2020    POPT
## 1 30372 -65649 -9.5 183821  26.5   1.4           1.6 227834
## 2 44299  9719  1.1 162109  18.8   NA           NA 429637

```

```

## 3    5190    3508      3.9  16352   18.3     0.2      0.2  63582
## 4   14486     470      0.2  55680   19.6     NA      NA  86704
## 5   56276 -19089     -1.8 239489   22.4     2.1      2.4 338124
## 6   58507 -12546     -1.1 241246   20.9     2.2      2.6 423973
##   POPTRT           geometry
## 1  32.9 MULTIPOLYGON (((22.97854 43...
## 2  49.7 MULTIPOLYGON (((8.56741 47....
## 3  71.3 MULTIPOLYGON (((33.77666 34...
## 4  30.6 MULTIPOLYGON (((19.77166 42...
## 5  31.6 MULTIPOLYGON (((14.48801 51...
## 6  36.7 MULTIPOLYGON (((5.0811 51.4...

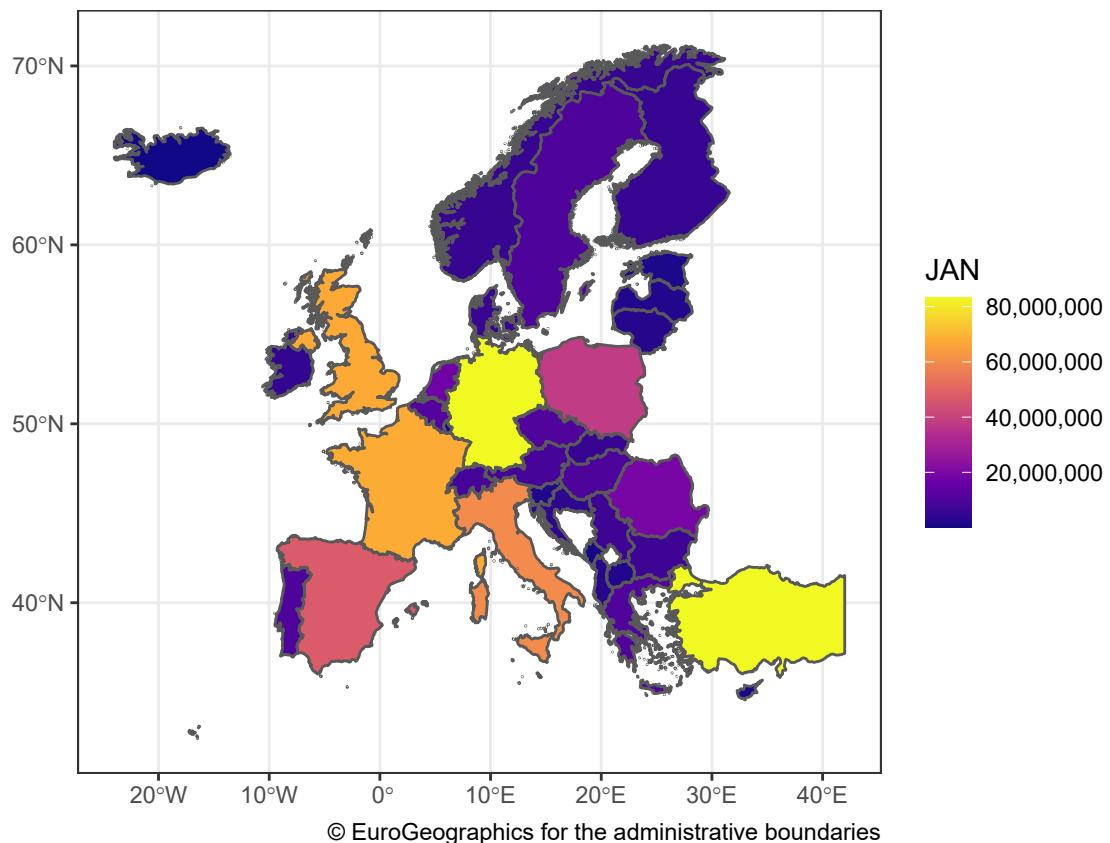
```

And now we can use the data in the aesthetics mapping to colour the countries.

```

ggplot(demomap) + aes(fill = JAN) + geom_sf() +
  theme_bw() +
  scale_fill_viridis_c(option = "plasma", label = scales::comma) +
  labs(caption = "@ EuroGeographics for the administrative boundaries",
       x = NULL, y = NULL)

```



Remember, we can improve the map display by using appropriate map projections. Here, we use the Lambert Conformal Conic Europe projection (EPSG 3034).

```

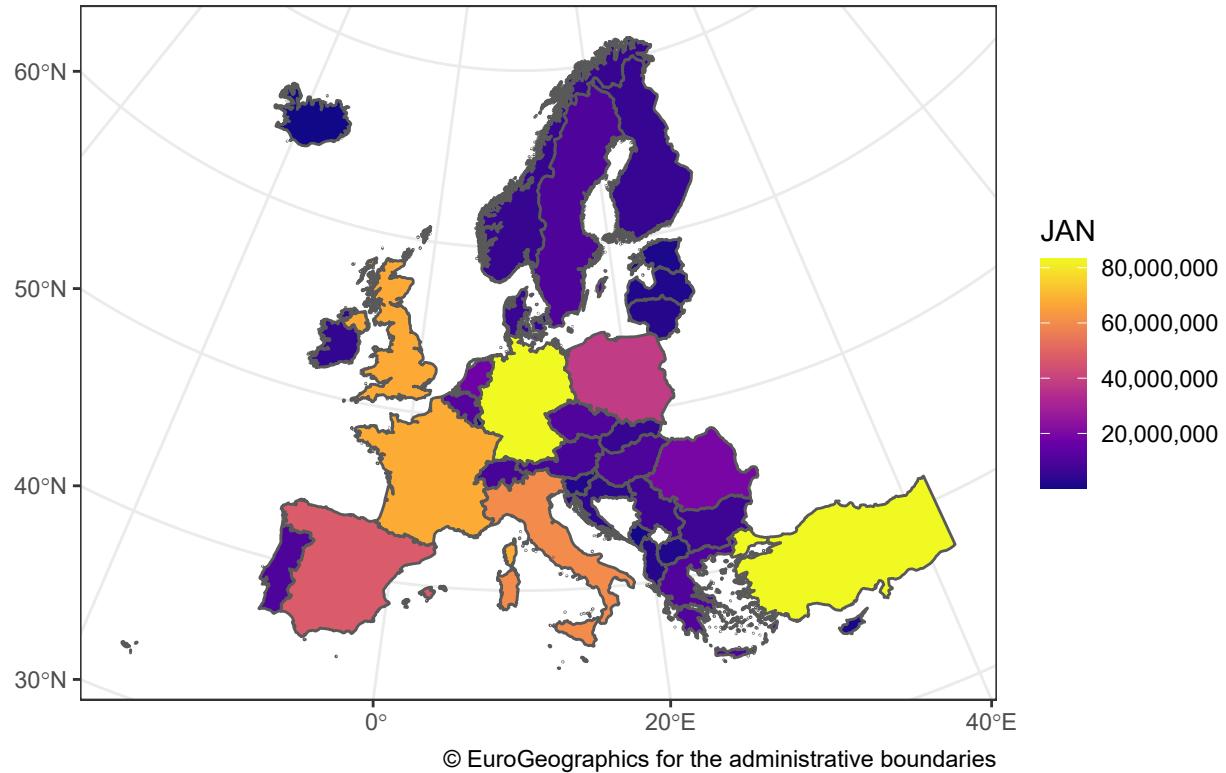
ggplot(demomap) + aes(fill = JAN) + geom_sf() +
  theme_bw() +
  coord_sf(crs = st_crs(3034)) +

```

```

scale_fill_viridis_c(option = "plasma", label = scales::comma) +
  labs(caption = "© EuroGeographics for the administrative boundaries",
       x = NULL, y = NULL)

```



The map above is a choropleth, which means we can also plot it as a distorted cartogram. Note that we first need to transform the map to use a Europe-specific CRS, in this case EPSG:23038.

```

library(cartogram)

## Warning: package 'cartogram' was built under R version 4.2.2

europeCartogram <- cartogram_cont(
  st_transform(demomap, crs = st_crs(3034)),
  weight = "JAN",
  itermax = 3 #limit to 3 iterations, so it's a bit faster
)

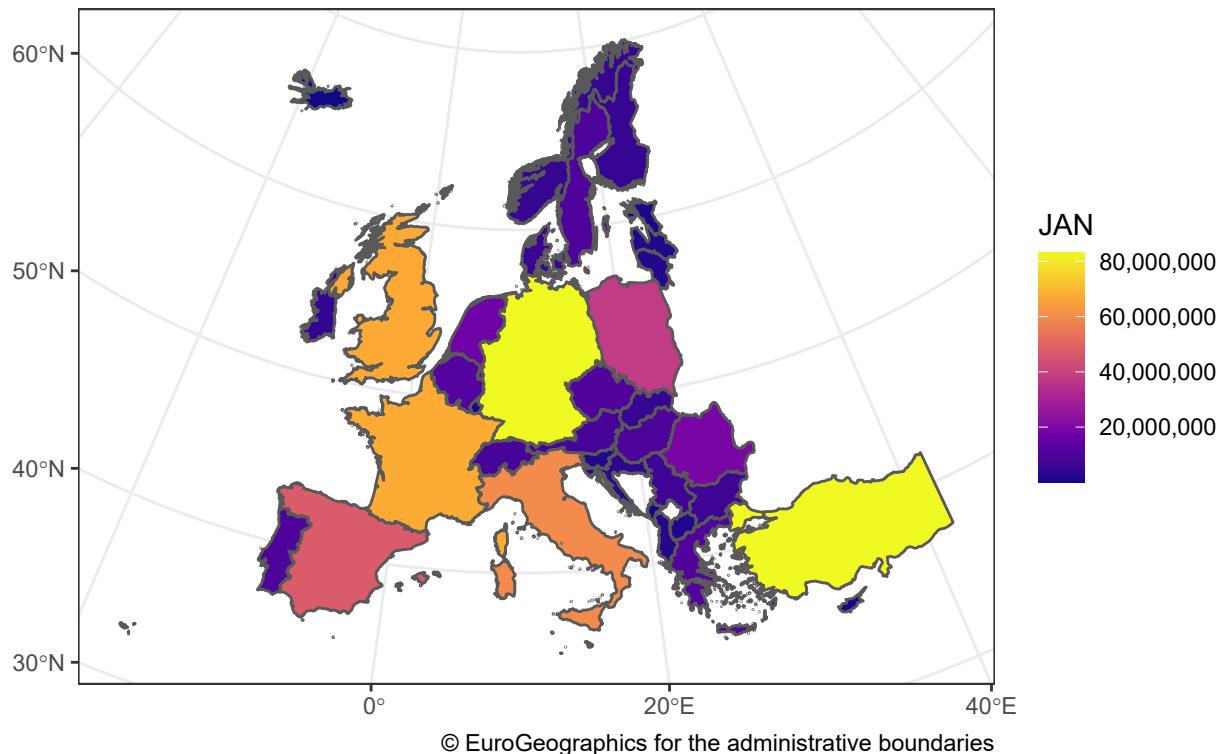
## Mean size error for iteration 1: 2.32998444936206

## Mean size error for iteration 2: 1.90279687313565

## Mean size error for iteration 3: 1.61516171385541

```

```
ggplot(europeCartogram) + aes(fill = JAN) + geom_sf() +
  theme_bw() +
  coord_sf(crs = st_crs(3034)) +
  scale_fill_viridis_c(option = "plasma", label = scales::comma) +
  labs(caption = "@ EuroGeographics for the administrative boundaries",
       x = NULL, y = NULL)
```



United Kingdom

You will often need to import maps from other sources. Mapping data is often provided in so-called *Shapefiles*. United Kingdom upper tier local authority (i.e. counties, unitary authorities, London boroughs, etc.) boundaries are available to download from geoportal.statistics.gov.uk. Note that there are different resolutions of the data. It's often enough to use generalised or even ultra generalised data, as these are much smaller and graphics render much quicker.

The *sf* package can read data in the *Shapefile* format, using the `st_read` function.

```
utla.map <- st_read(
  "UTLA2019/Counties_and_Unitary_Authorities_(December_2019)_Boundaries_UK_BUC.shp")
```

```
## Reading layer `Counties_and_Unitary_Authorities_(December_2019)_Boundaries_UK_BUC` from data source
##   using driver 'ESRI Shapefile'
## Simple feature collection with 216 features and 10 fields
```

```

## Geometry type: MULTIPOLYGON
## Dimension: XY
## Bounding box: xmin: -116.1928 ymin: 7054.1 xmax: 655653.8 ymax: 1218618
## Projected CRS: OSGB36 / British National Grid

```

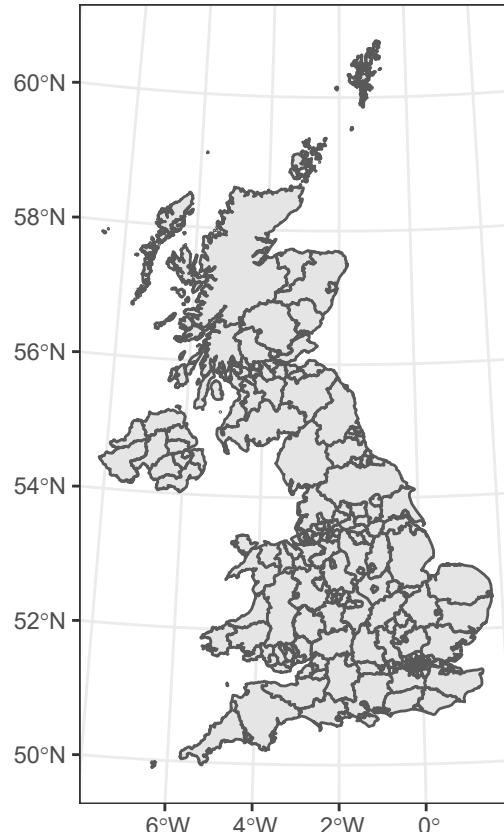
```
head(as.data.frame(utla.map))
```

```

##   objectid ctyua19cd      ctyua19nm ctyua19nmw bng_e bng_n      long
## 1       1 E06000001    Hartlepool     <NA> 447160 531474 -1.27018
## 2       2 E06000002 Middlesbrough     <NA> 451141 516887 -1.21099
## 3       3 E06000003 Redcar and Cleveland     <NA> 464361 519597 -1.00608
## 4       4 E06000004 Stockton-on-Tees     <NA> 444940 518183 -1.30664
## 5       5 E06000005 Darlington     <NA> 428029 515648 -1.56835
## 6       6 E06000006 Halton     <NA> 354246 382146 -2.68853
##   lat st_areasha st_lengths           geometry
## 1 54.67614    96845510  50305.33 MULTIPOLYGON (((448986 5367...
## 2 54.54467    52908459  34964.41 MULTIPOLYGON (((451752.7 52...
## 3 54.56752   248679106  83939.75 MULTIPOLYGON (((451965.6 52...
## 4 54.55691   207159134  87075.86 MULTIPOLYGON (((451965.6 52...
## 5 54.53534   198812771  91926.84 MULTIPOLYGON (((419709.3 51...
## 6 53.33424    80285488  66618.91 MULTIPOLYGON (((355156.2 38...

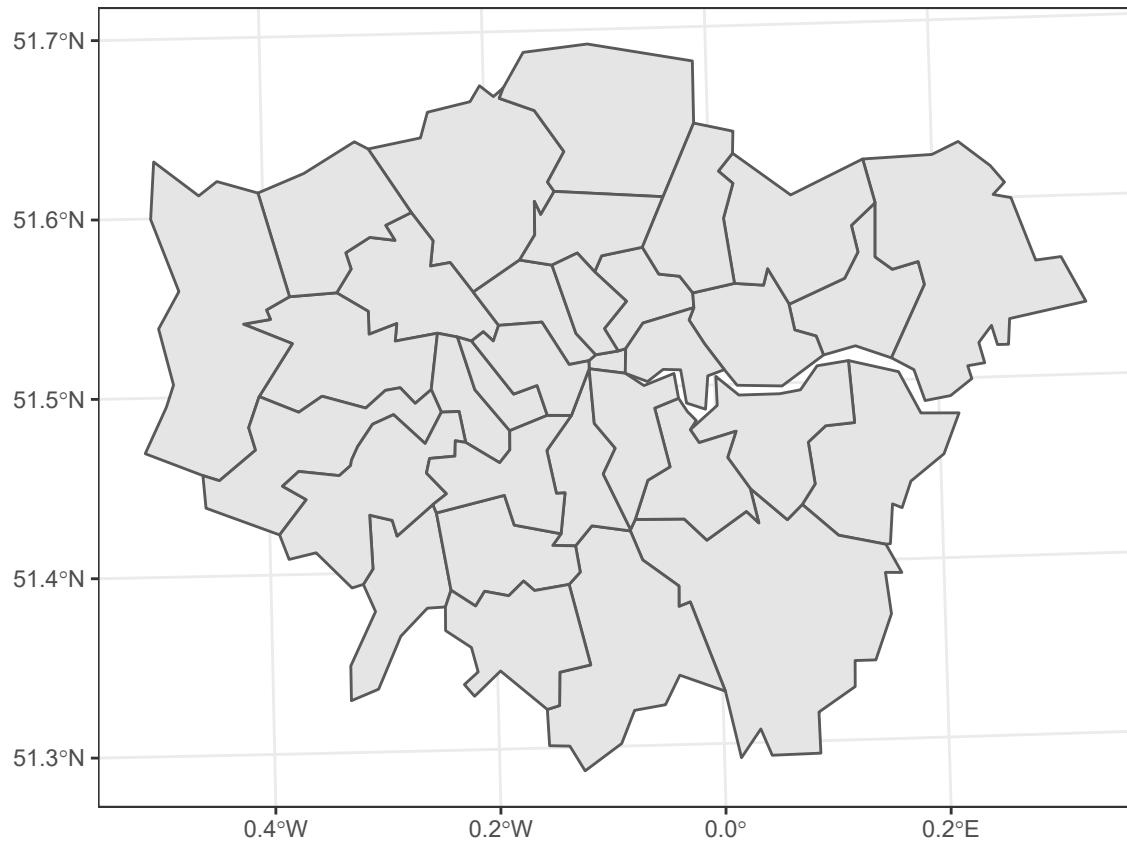
```

```
ggplot(utla.map) + geom_sf() + theme_bw()
```



If we only want to display a part of the map we need to filter the data. Here, we create a map of the London boroughs (all London boroughs start with E09).

```
ldn.map <- utla.map[substr(utla.map$ctyua19cd, 1, 3) == "E09",]
ggplot(ldn.map) + geom_sf() +
  theme_bw()
```



You might notice that the map is somewhat tilted. This is because the imported shapefile uses the British National Grid coordinate reference system (CRS).

```
st_crs(ldn.map)

## Coordinate Reference System:
##   User input: OSGB36 / British National Grid
##   wkt:
##     PROJCRS["OSGB36 / British National Grid",
##             BASEGEOGCRS["OSGB36",
##                         DATUM["Ordnance Survey of Great Britain 1936",
##                                ELLIPSOID["Airy 1830",6377563.396,299.3249646,
##                                         LENGTHUNIT["metre",1]]],
##                         PRIMEM["Greenwich",0,
##                                ANGLEUNIT["degree",0.0174532925199433]],
##                         ID["EPSG",4277]],
##             CONVERSION["British National Grid",
##                     METHOD["Transverse Mercator",
##                            ID["EPSG",9807]],
##                     PARAMETER["Latitude of natural origin",49,
##                               ANGLEUNIT["degree",0.0174532925199433],
```

```

##           ID["EPSG",8801]],
## PARAMETER["Longitude of natural origin",-2,
##           ANGLEUNIT["degree",0.0174532925199433],
##           ID["EPSG",8802]],
## PARAMETER["Scale factor at natural origin",0.9996012717,
##           SCALEUNIT["unity",1],
##           ID["EPSG",8805]],
## PARAMETER["False easting",400000,
##           LENGTHUNIT["metre",1],
##           ID["EPSG",8806]],
## PARAMETER["False northing",-100000,
##           LENGTHUNIT["metre",1],
##           ID["EPSG",8807]],
## CS[Cartesian,2],
##     AXIS["(E)",east,
##         ORDER[1],
##         LENGTHUNIT["metre",1]],
##     AXIS["(N)",north,
##         ORDER[2],
##         LENGTHUNIT["metre",1]],
## USAGE[
##     SCOPE["Engineering survey, topographic mapping."],
##     AREA["United Kingdom (UK) - offshore to boundary of UKCS within 49°45'N to 61°N and 9°W to 2
##          BBOX[49.75,-9,61.01,2.01]],
##     ID["EPSG",27700]]

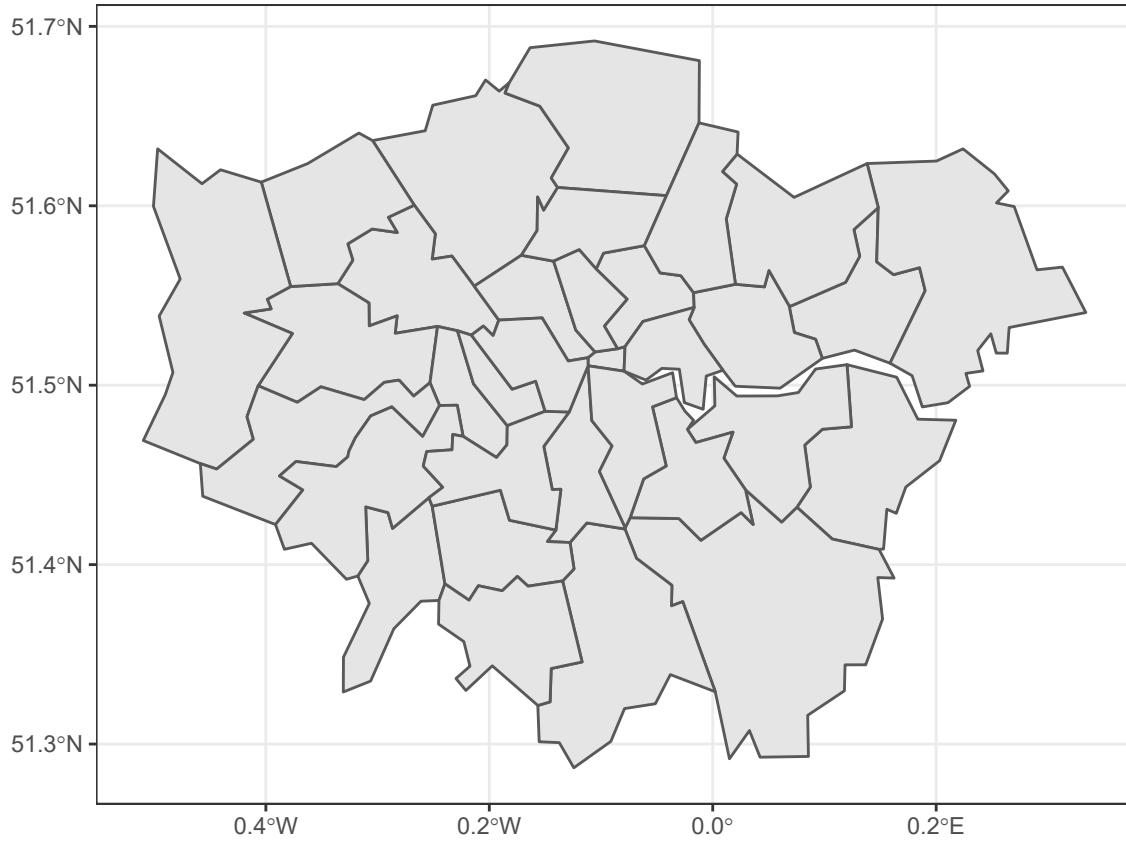
```

We can transform the map to use *WGS84* CRS.

```

ggplot(ldn.map) + geom_sf() +
  theme_bw() +
  coord_sf(crs = st_crs(4326))

```



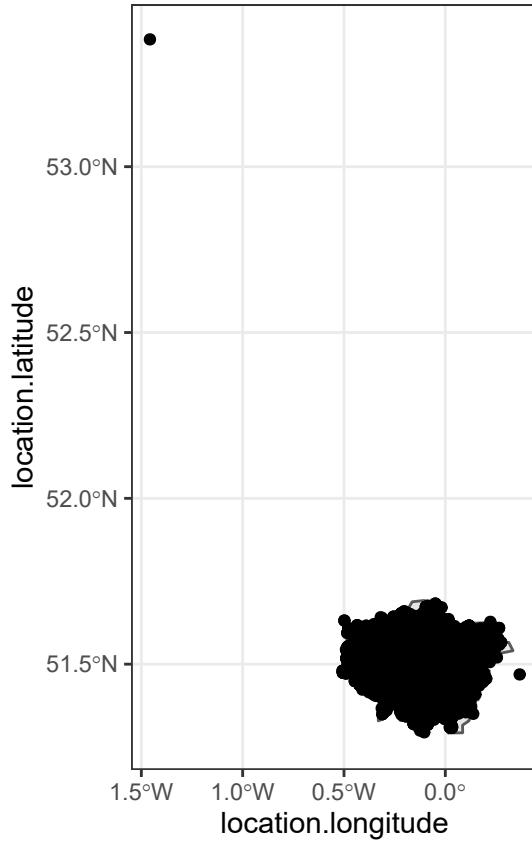
Now let's add some data to the map. Let's display the locations of police stop and searches on the map.

We can load the data from [data.police.uk](https://data.police.uk/api/stops-force?force=metropolitan&date), as a JSON object with `readJSON`. The `flatten` method converts this into a usable data frame. Then we need to make sure that the coordinates are actually stored as numbers, and finally we remove any records without coordinates.

```
library(jsonlite)
#stop_search <- fromJSON("https://data.police.uk/api/stops-force?force=metropolitan&date")
stop_search <- fromJSON("stops-force.json")
stop_search <- flatten(stop_search)
stop_search$location.latitude <-
  as.numeric(stop_search$location.latitude)
stop_search$location.longitude <-
  as.numeric(stop_search$location.longitude)
stop_search <- stop_search[!is.na(stop_search$location.latitude),]
```

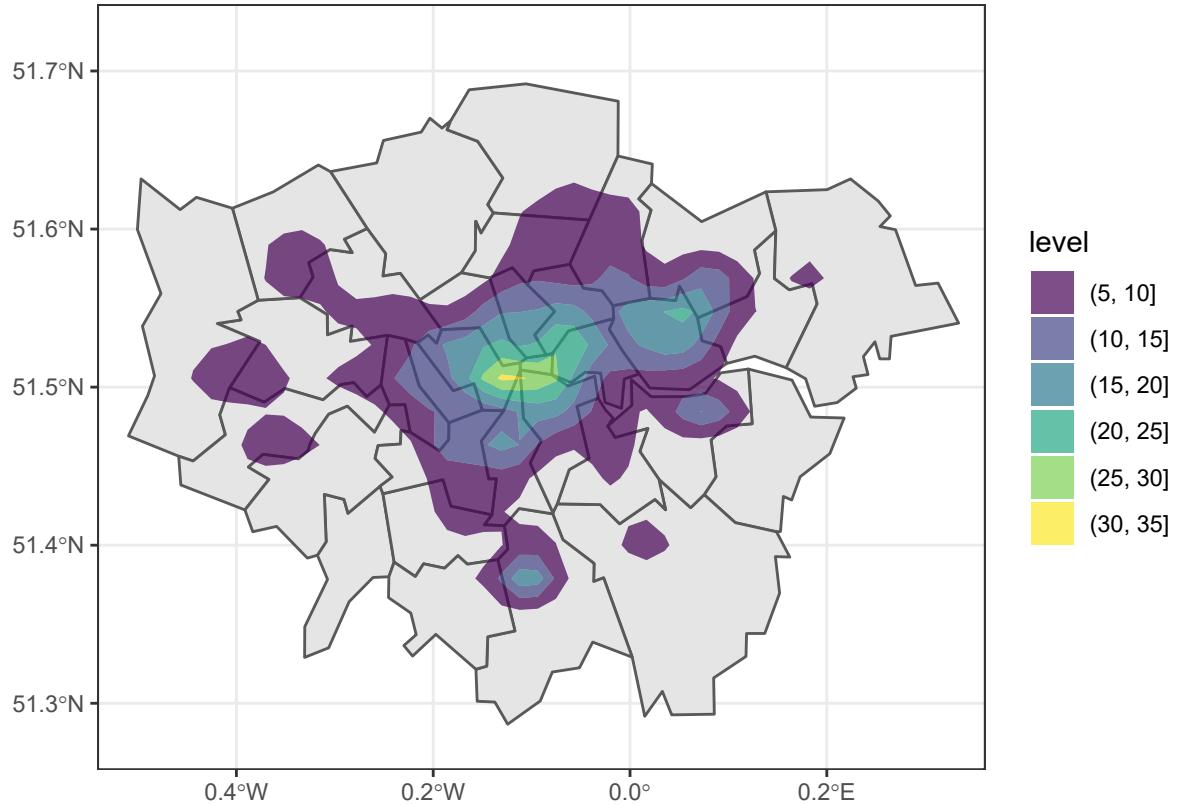
We use `geom_point()` to overlay the stop and searches on the map.

```
ggplot(ldn.map) + geom_sf() +
  geom_point(data = stop_search,
             aes(location.longitude, location.latitude)) +
  theme_bw() +
  coord_sf(crs = st_crs(4326))
```



That is clearly not quite right! There are way too many points on the map and two incidents are far outside of London. By creating a density map we show the density of incidents at a certain location instead of the actual dots and by specifying `xlim` and `ylim` we make sure to only show London.

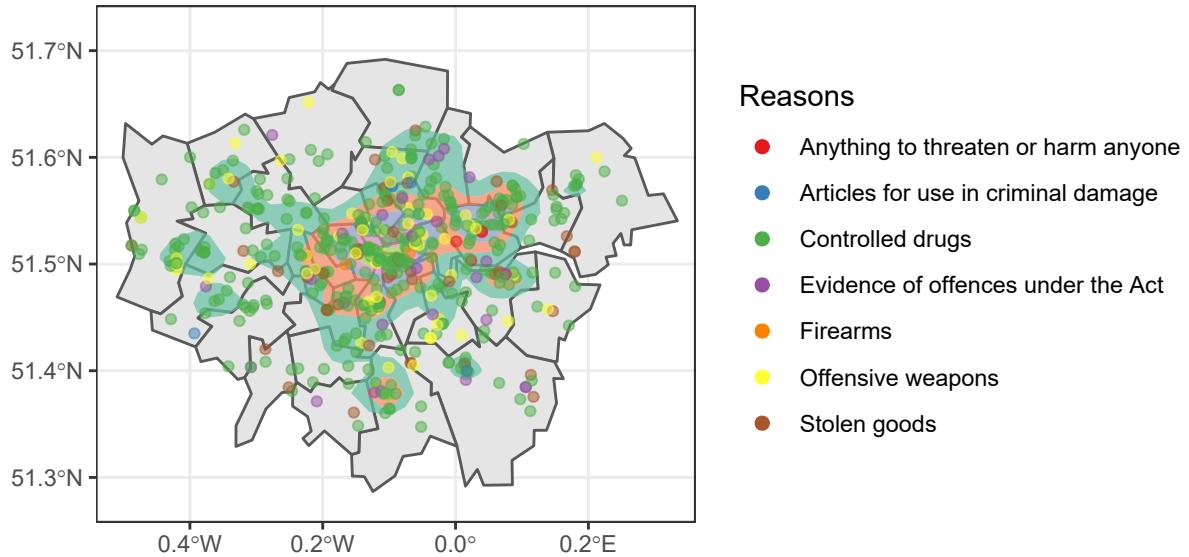
```
ggplot(ldn.map) + geom_sf() +
  geom_density2d_filled(data = stop_search, alpha = 0.7,
                        breaks = seq(5, 35, 5),
                        aes(location.longitude, location.latitude)) +
  theme_bw() +
  coord_sf(crs = st_crs(4326), xlim = c(-0.5, 0.32), ylim = c(51.28, 51.72)) +
  labs(x= "", y = "")
```



Finally, let's also display a sample of the stop and searches to get an idea for the reasons of the stop and searches.

```
stop_search_sample <- stop_search[sample(nrow(stop_search), 500),]
stop_search_sample <-
  stop_search_sample[!is.na(stop_search_sample$object_of_search),]
ggplot(ldn.map) + geom_sf() +
  geom_density2d_filled(data = stop_search, alpha = 0.7,
                        breaks = seq(5, 35, 5),
                        aes(location.longitude, location.latitude)) +
  geom_point(data = stop_search_sample, alpha = 0.5,
             aes(location.longitude, location.latitude,
                 colour = object_of_search)) +
  theme_bw() +
  scale_fill_brewer(palette = "Set2") +
  scale_color_brewer(palette = "Set1") +
  coord_sf(crs = st_crs(4326), xlim = c(-0.5, 0.32), ylim = c(51.28, 51.72)) +
  guides(fill = "none",
         colour = guide_legend(
           title = "Reasons",
           override.aes = list(
             alpha = 1,
             size = 2,
             fill = NA,
             linetype = 0
           )))
```

```
labs(x= "", y = "")
```



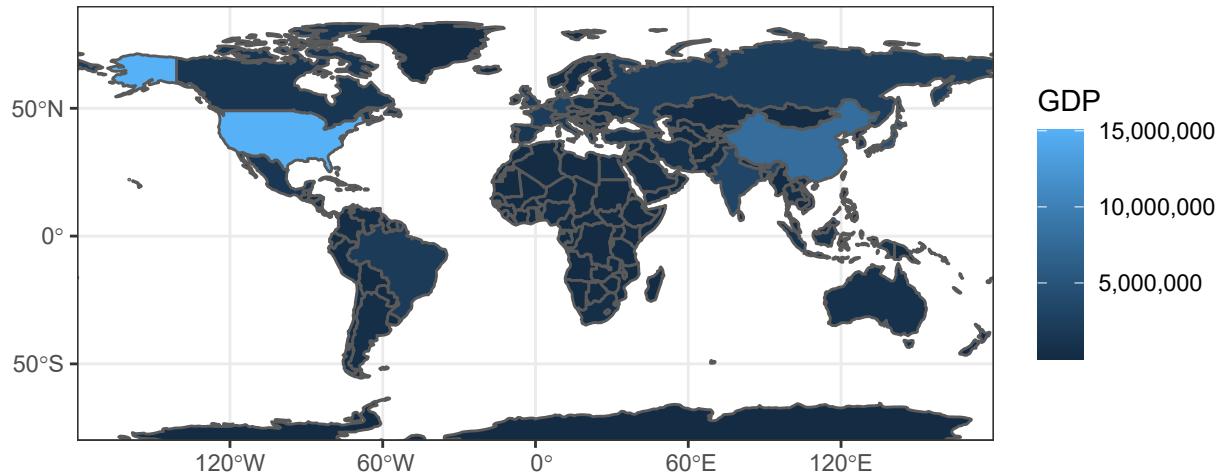
Exercises

1. Create a world map showing the GDP per country. Note, the object returned by `ne_countries()` contains the estimated GDP in the column `gdp_md_est`. Should you use a linear or an exponential scale?
2. Create a world map showing which economic group each country belongs to. Use the `economy` column.
3. Using the `ne_countries()` function, create a map showing only Africa. Show the population of each country.
4. The eurostat statistic [tgs00047](#) contains the percentage of households that have internet access at home by NUTS level 2 regions. Visualise the newest available data.
5. The file `ukcities.csv` contains a list of all cities in the UK, their coordinates and population. Draw a map of the UK and add the cities as dots with the size proportional to the cities' population. Bonus: Add labels for the five largest cities.

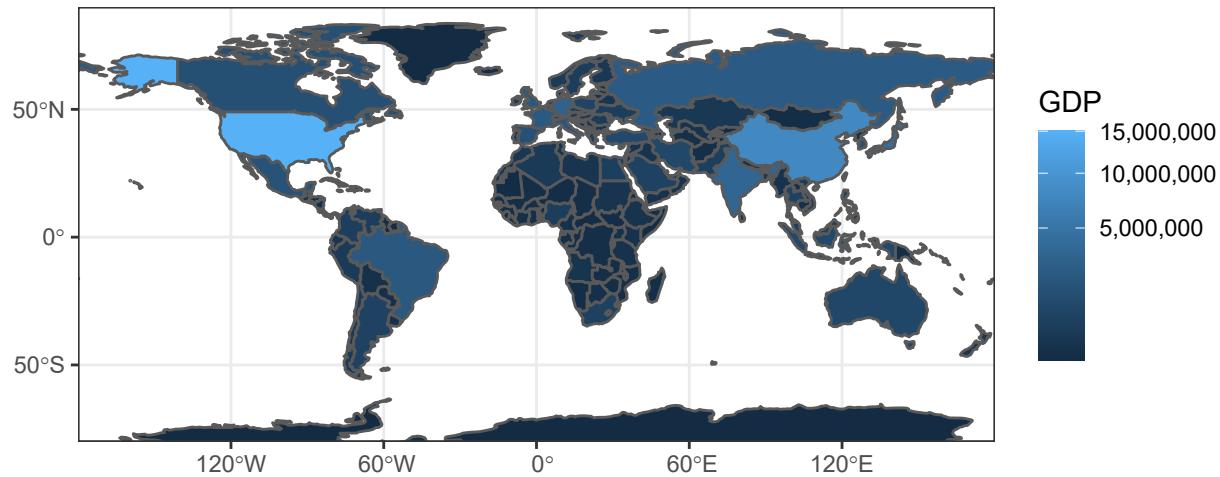
Solutions

1. Create a world map showing the GDP per country. Note, the object returned by `ne_countries()` contains the estimated GDP in the column `gdp_md_est`. Should you use a linear or an exponential scale?

```
world <- ne_countries(returnclass = "sf")
ggplot(world) + geom_sf() + aes(fill = gdp_md_est) +
  scale_fill_continuous(label = scales::comma) +
  labs(fill = "GDP") +
  theme_bw() + coord_sf(ylim = c(-80, 90), expand = F)
```

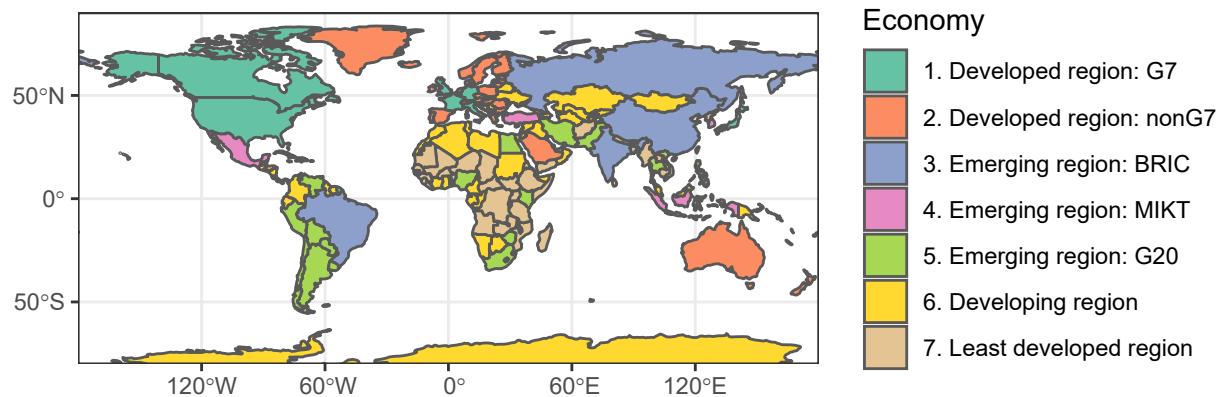


```
ggplot(world) + geom_sf() + aes(fill = gdp_md_est) +
  scale_fill_continuous(label = scales::comma, trans = "sqrt") +
  labs(fill = "GDP") +
  theme_bw() + coord_sf(ylim = c(-80, 90), expand = F)
```



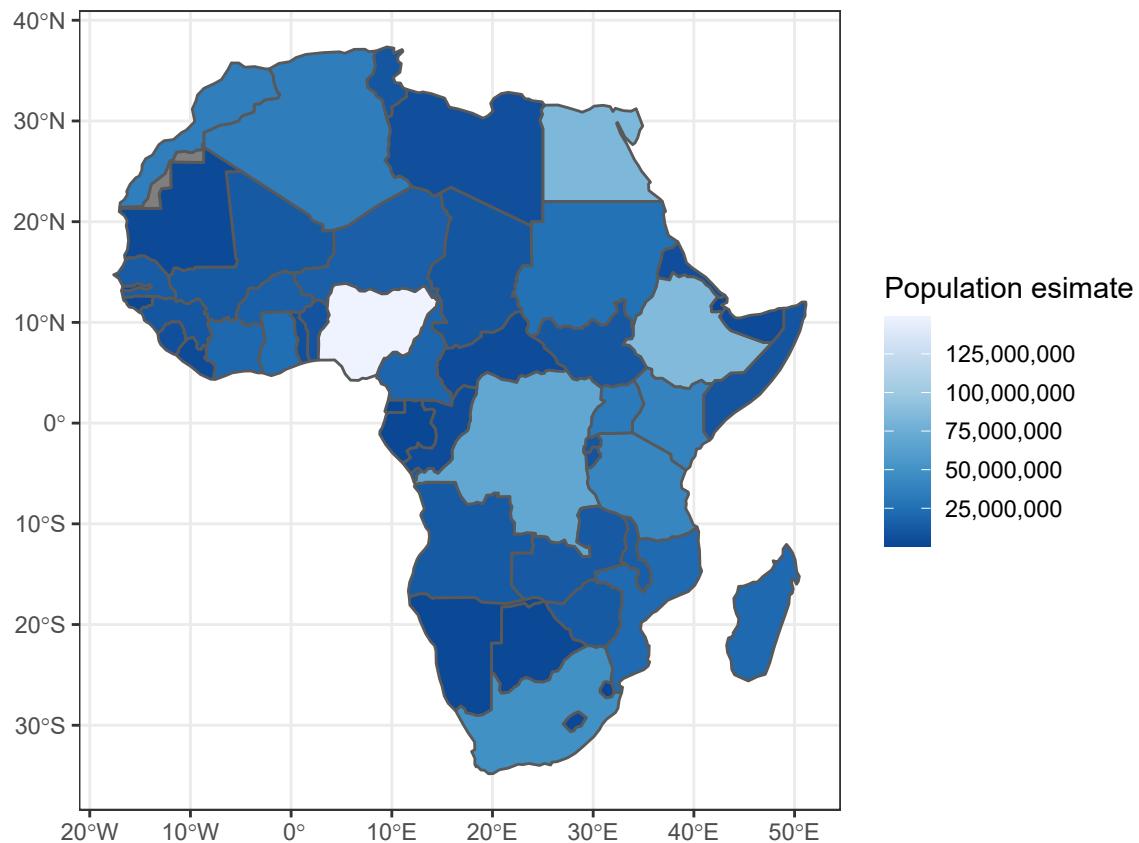
2. Create a world map showing which economic group each county belongs to. Use the `economy` column.

```
world <- ne_countries(returnclass = "sf")
ggplot(world) + geom_sf() + aes(fill = economy) +
  scale_fill_brewer(palette = "Set2") +
  labs(fill = "Economy") +
  theme_bw() + coord_sf(ylim = c(-80, 90), expand = F)
```



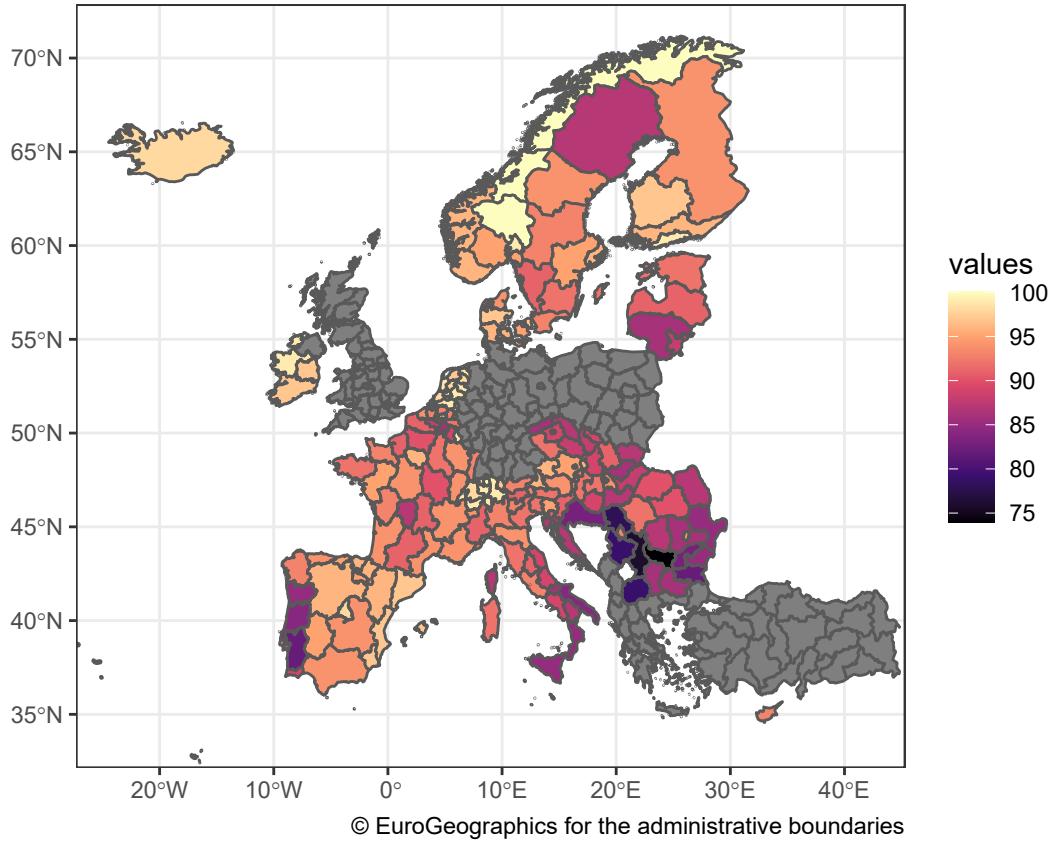
3. Using the `ne_countries()` function, create a map showing only Africa. Show the population of each country.

```
africa <- ne_countries(returnclass = "sf", continent = "Africa")
ggplot(africa) + geom_sf() + aes(fill = pop_est) +
  scale_fill_distiller(label = scales::comma, breaks = seq(0, 200e+6, 25e+6)) +
  labs(fill = "Population estimate") +
  theme_bw()
```



4. The eurostat statistic `tgs00047` contains the percentage of households that have internet access at home by NUTS level 2 regions. Visualise the newest available data.

```
euro.map.regional <- get_eurostat_geospatial(nuts = 2)
data <- get_eurostat("tgs00047")
data <-
  data %>%
  group_by(geo) %>%
  slice_max(order_by = time, n = 1)
data.map <- left_join(euro.map.regional, data, by = c("id" = "geo"))
ggplot(data.map) + geom_sf() + aes(fill = values) +
  theme_bw() +
  scale_fill_viridis_c(option = "magma") +
  coord_sf(xlim=c(-24,42), ylim=c(34, 71)) +
  labs(caption = "@ EuroGeographics for the administrative boundaries",
       x = NULL, y = NULL)
```



5. The file `ukcities.csv` contains a list of all cities in the UK, their coordinates and population. Draw a map of the UK and add the cities as dots with the size proportional to the cities' population. Bonus: Add labels for the five largest cities.

```
library(readr)
ukcities <- read_csv("ukcities.csv")

## # Rows: 51 Columns: 4
## -- Column specification -----
## Delimiter: ","
## chr (1): city
## dbl (3): lat, lng, population
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

utla.map <- st_read("UTLA2019/Counties_and_UnitaryAuthorities_(December_2019)_Boundaries_UK_BUC.shp")

## # Reading layer `Counties_and_UnitaryAuthorities_(December_2019)_Boundaries_UK_BUC` from data source
##   using driver 'ESRI Shapefile'
## Simple feature collection with 216 features and 10 fields
## Geometry type: MULTIPOLYGON
## Dimension:      XY
## Bounding box:  xmin: -116.1928 ymin: 7054.1 xmax: 655653.8 ymax: 1218618
## Projected CRS: OSGB36 / British National Grid
```

```

largest <- slice_max(ukcities, order_by = population, n = 5)

ggplot(utla.map) + geom_sf() +
  geom_point(data = ukcities, aes(lng, lat, size = population), colour = "orchid3") +
  geom_text(data = largest, aes(lng, lat, label = city),
            size = 3, fontface = "bold") +
  scale_size_continuous(label = scales::comma,
                        guide = guide_legend(reverse=TRUE)) +
  theme_bw() +
  coord_sf(crs = st_crs(4326)) +
  labs(x = "", y = "")

```

